

HEALTH AI INTELLIGENT HEALTH CARE ASSISTANT PROJECT DOCUMENTATION

1. Introduction

Project Title: Health AI Intelligent Health Care Assistant

Team member : MUTHU GANAPATHI R P

Team member : HAJA MOHAIDEEN B

Team member : AAKASH M

Team member : MOHAMED MUFEETH J

Team member : DONY A

2. Project Overview

Purpose: The Medical AI Assistant is designed to provide informational support on possible

Features:

Disease Prediction: Users can enter symptoms and receive suggestions for possible medical

Treatment Plan Generation: The assistant can generate personalized treatment plans based on

User Interface: The application uses a Gradio interface with separate tabs for each function.

LLM Integration: It uses the `ibm-granite/granite-3.2-2b-instruct` model for generating responses.

3. Solution Components

AI Model: The project uses `ibm-granite/granite-3.2-2b-instruct` for natural language understanding.

Frontend: The frontend is a web interface built with the Gradio library. It includes text boxes for

Backend: The logic is handled by a Python script that uses the Hugging Face Transformers

4. Technology Stack

Frameworks: Gradio

Libraries: `torch` and `transformers` are used for loading and running the AI model.

API Endpoints: The code itself does not define explicit API endpoints, but the Gradio interface

5. Prompt Engineering

Purpose: Prompts are carefully crafted to guide the AI model to provide medical information. The `disease_prediction` and `treatment_plan` functions use specific prompts that include the user's input and a clear disclaimer about consulting a healthcare professional.

6. Deployment

The application can be launched and shared via a public URL using the `app.launch(share=True)` command. The provided document also shows that the deployment was in a Colab notebook.

7. API Endpoints

The application's functionality is handled by two main Python functions: `disease_prediction` and `treatment_plan`. These functions take user inputs and return AI-generated responses.

8. Authentication

The provided code does not include authentication. A warning is noted that the Hugging Face Hub token does not exist, but authentication is optional to access public models.

9. User Interface

The interface is created using `gr.Blocks()`. It features a main title and a disclaimer using `gr.Markdown` and is organized into two tabs with `gr.Tabs` and `gr.TabItem` for different functionalities. The layout uses `gr.Row` and `gr.Column` to arrange input and output components.

10. Testing

- **Unit Testing**– Prompt responses and ML models.
- **API Testing**– Swagger & Postman.
- **Manual Testing**– File uploads, summarization, anomaly detection.
- **Edge Cases**– Invalid inputs, empty PDFs, missing API keys.

11. Screenshot

```
[ ] | pip install transformers torch  
| gradio -q
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-packages:  
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages (2.8.1  
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (fi  
Requirement already satisfied: huggingface-hub<1.0,>=0.34.0 in /usr/local/lib/python3  
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-packages  
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-pack  
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages  
Requirement already satisfied: regex<2019.12.17 in /usr/local/lib/python3.12/dist-pa  
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages (fi  
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/python3.  
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.12/dist-pi  
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.12/dist-packages  
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12.  
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages  
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packag  
Requirement already satisfied: networkx in /usr/local/lib/python3.12/dist-packages (fi  
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.12/dist-packages (fron  
Requirement already satisfied: fsspec in /usr/local/lib/python3.12/dist-packages (fron  
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/pytho  
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/py  
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/pytho  
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3  
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/local/lib/python3  
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.  
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python:  
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/pythor  
Requirement already satisfied: nvidia-cuspars-cu12==12.5.4.2 in /usr/local/lib/pythor
```



```
HEALTH.AI
+ <> + T
Connect T4
/usr/local/lib/python3.12/dist-packages/huggingface_hub
The secret 'HF_TOKEN' does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (http
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public m
warnings.warn(
tokenizer_config.json: 8.88k/? [00:00<00:00, 256kB/s]
vocab.json: 777k/? [00:00<00:00, 6.02MB/s]
merges.txt: 442k/? [00:00<00:00, 8.93MB/s]
tokenizer.json: 3.48M/? [00:00<00:00, 47.2MB/s]
added_tokens.json: 100% 87.0/87.0 [00:00<00:00, 3.19kB/s]
special_tokens_map.json: 100% 701/701 [00:00<00:00, 21.6kB/s]
config.json: 100% 786/786 [00:00<00:00, 22.5kB/s]
'torch_dtype' is deprecated! Use 'dtype' instead!
model.safetensors.index.json: 29.8k/? [00:00<00:00, 3.11MB/s]
Fetching 2 files: 100% 2/2 [01:11<00:00, 71.53s/it]
model-00001-of-5.00G/5.00G [01:11<00:00, 76.2MB/s]
00002.safetensors: 100%
model-00002-of-67.1M/67.1M [00:01<00:00, 53.3MB/s]
00002.safetensors: 100%
Loading checkpoint shards: 100% 2/2 [00:19<00:00, 8.11s/it]
generation_config.json: 100% 137/137 [00:00<00:00, 3.77kB/s]
Colab notebook detected. To show errors in colab notebook, set debug=True in launch(
* Running on public URL: https://095cdbd76ff5073b14.gradio.live
This share link expires in 1 week. For free permanent hosting and GPU upgrades, run
```



```
import gradio as gr
import torch
from transformers import AutoTokenizer, AutoModelForCausalLM

# Load model and tokenizer
model_name = "ibm-granite/granite-3.2-2b-instruct"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    torch_dtype=torch.float16 if torch.cuda.is_available() else torch.float32,
    device_map="auto" if torch.cuda.is_available() else None
)

if tokenizer.pad_token is None:
    tokenizer.pad_token = tokenizer.eos_token

def generate_response(prompt, max_length=1024):
    inputs = tokenizer(prompt, return_tensors="pt", truncation=True, max_length=512)

    if torch.cuda.is_available():
        inputs = {k: v.to(model.device) for k, v in inputs.items()}

    with torch.no_grad():
        outputs = model.generate(
            **inputs,
            max_length=max_length,
            temperature=0.7,
            do_sample=True,
            pad_token_id=tokenizer.eos_token_id
        )

    response = tokenizer.decode(outputs[0], skip_special_tokens=True)
    response = response.replace(prompt, "").strip()
    return response

def disease_prediction(symptoms):
    prompt = f"Based on the following symptoms, provide possible medical conditions and generate a treatment plan."
    return generate_response(prompt, max_length=1200)

def treatment_plan(condition, age, gender, medical_history):
    prompt = f"Generate personalized treatment suggestions for the following patient information: {condition}, {age}, {gender}, {medical_history}."
    return generate_response(prompt, max_length=1200)

# Create Gradio interface
with gr.Blocks() as app:
    gr.Markdown("# Medical AI Assistant")
```



```
gr.Markdown("**Disclaimer: This is for informational  
with gr.Tabs():  
    with gr.TabItem("Disease Prediction"):  
        with gr.Row():  
            with gr.Column():  
                symptoms_input = gr.Textbox(  
                    label="Enter Symptoms",  
                    placeholder="e.g., fever, headache, cough, fatigue...",  
                    lines=4  
                )  
                predict_btn = gr.Button("Analyze Symptoms")  
  
            with gr.Column():  
                prediction_output = gr.Textbox(label="Possible Conditions & Recommendations")  
  
        predict_btn.click(disease_prediction, inputs=symptoms_input, outputs=prediction_output)  
  
    with gr.TabItem("Treatment Plans"):  
        with gr.Row():  
            with gr.Column():  
                condition_input = gr.Textbox(  
                    label="Medical Condition",  
                    placeholder="e.g., diabetes, hypertension, migraine...",  
                    lines=2  
                )  
                age_input = gr.Number(label="Age", value=30)  
                gender_input = gr.Dropdown(  
                    choices=["Male", "Female", "Other"],  
                    label="Gender",  
                    value="Male"  
                )  
                history_input = gr.Textbox(  
                    label="Medical History",  
                    placeholder="Previous conditions, allergies, medications or None",  
                    lines=3  
                )  
                plan_btn = gr.Button("Generate Treatment Plan")  
  
            with gr.Column():  
                plan_output = gr.Textbox(label="Personalized Treatment Plan", lines=20)  
  
        plan_btn.click(treatment_plan, inputs=[condition_input, age_input, gender_input, history_input], outputs=plan_output)  
app.launch(share=True)
```

Colab notebook detected. To show errors in colab notel ↑ ↓ ✦ 🔗 ⚙️ 📄 🗑️ ⋮

* Running on public URL: <https://44e8be4f8de382703a.g...>



This share link expires in 1 week. For free permanent hosting and GPU upgrades, run

Medical AI Assistant

Disclaimer: This is for informational purposes only. Always consult healthcare professionals for medical advice.

Disease Prediction

Treatment Plans

Enter Symptoms

e.g., fever, headache, cough, fatigue...

Analyze Symptoms

Possible Conditions & Recommendations

12. Known Issues

- Occasional long response time for large PDFs.
- Forecasting limited to structured CSV data.
- Requires stable internet for BMAP access.

13. Future Enhancements

- Add voice-based interaction.
- Expand forecasting to include traffic & pollution data.
- Develop a mobile app version.
- Integrate with IoT smart sensors.
- Support multi-language outputs for local communities.